

Determining Synchronous vs. Asynchronous Design

A Decision Framework for Designing Business Capability Interactions

Author: Placida Fernando


Purpose

When a business capability is exposed as a service interaction, one of the first architectural decisions is whether that interaction should be synchronous (request/response, caller blocks and waits) or asynchronous (fire-and-forget, event-driven, or message-based, caller continues without waiting). This decision affects consumer experience, system coupling, resiliency, and scalability. This paper provides a concise decision framework enterprise architects and solution designers can apply consistently when defining integration patterns for a business capability.


Core Definitions

- Synchronous: The requester sends a request and blocks/waits until the provider returns a response before proceeding. Tight temporal coupling exists between the two parties.
- Asynchronous: The requester sends a request (or publishes an event) and continues processing independently. The response, if any, is delivered later via callback, polling, event, or message queue.

Synchronous Pattern



Asynchronous Pattern



Decision Criteria

Evaluate the capability interaction against the criteria below. The more conditions that point to a single column, the clearer the design choice; mixed results indicate a hybrid pattern (e.g., synchronous acknowledgment with asynchronous completion) may be appropriate.

Decision Criterion	Favors Synchronous	Favors Asynchronous
Response dependency	Caller needs the result immediately to proceed (e.g., validate an order before confirming to a customer)	Caller can continue without waiting; result is consumed later or by another process
Consistency requirement	Strong/immediate consistency needed across the transaction boundary	Eventual consistency is acceptable
Processing duration	Operation completes quickly and predictably (typically sub-second to a few seconds)	Operation is long-running, batch-oriented, or has unpredictable duration
Coupling tolerance	Tight coupling between consumer and provider is acceptable for this interaction	Consumer and provider should evolve and scale independently
Failure handling	Immediate error feedback to the caller is required	Retries, dead-letter handling, and eventual delivery are acceptable
Load & scalability	Traffic is steady and provider capacity comfortably exceeds demand	Traffic is bursty, unpredictable, or provider capacity is constrained (needs buffering)
Business criticality of decoupling	A direct, synchronous dependency does not threaten availability of the wider system	Provider outages must not cascade to the consumer or the business process
Integration span	Interaction is within a bounded, tightly related domain (e.g., a single microservice call chain)	Interaction crosses domain/organizational boundaries or triggers multiple downstream capabilities

Recommended Approach

- Start from the business requirement, not the technology: identify whether the calling process genuinely cannot proceed without an immediate result.
- Score the interaction against each criterion above; a majority favoring one side indicates the natural design.
- Where criteria conflict (e.g., user needs speed but the operation is long-running), consider a hybrid: synchronous acceptance/validation plus asynchronous processing and notification (e.g., HTTP 202 Accepted with a webhook or polling endpoint).
- Re-validate the decision whenever volume, criticality, or downstream dependencies change materially, the right pattern at launch may not remain right at scale.

Worked Example: Order Placement Capability

Consider an e-commerce "Place Order" business capability, which internally triggers payment authorization, inventory reservation, invoicing, and shipment scheduling.

- Payment authorization: The customer is waiting on-screen and cannot complete checkout without a decision. Response dependency and consistency needs are high → design as synchronous.
- Inventory reservation: Needs to be consistent enough to prevent overselling but can complete in milliseconds → typically synchronous, within the same transaction boundary as payment.
- Invoice generation: The customer does not need to wait for a PDF invoice to be generated and emailed → design as asynchronous, triggered by an "OrderPlaced" event.

- Shipment scheduling: Depends on a third-party logistics provider with variable response times and its own outages → design as asynchronous to avoid cascading failures into the checkout path.

Result: the overall "Place Order" capability is a hybrid, a synchronous core (payment + inventory) that returns an immediate confirmation, followed by an asynchronous fan-out of downstream steps driven by a domain event.

Decision Tree Summary

- Does the caller need the result to proceed right now? No → Asynchronous. Yes → continue.
- Is the operation short, predictable, and low-risk of failure? No → Asynchronous (with sync acknowledgment if needed). Yes → continue.
- Must the consumer and provider remain available independently of one another? Yes, strict isolation required → Asynchronous. No → continue.
- None of the above rule it out → Synchronous is appropriate.

Related Architectural Patterns

The following established patterns are commonly paired with asynchronous designs identified through this framework:

- **Saga pattern:** Coordinates a sequence of local transactions across services using asynchronous events or commands, with compensating actions on failure. Used when a business process (like order placement) spans multiple services that cannot share a single ACID transaction.
- **Event sourcing:** Persists state changes as an immutable sequence of events rather than just current state, allowing asynchronous consumers to rebuild views or trigger downstream processing from the event log.
- **CQRS (Command Query Responsibility Segregation):** Separates write (command) and read (query) models, often using asynchronous event propagation to keep read models eventually consistent with the write side, useful when synchronous reads would otherwise create tight coupling to the write path.

These patterns are not required for every asynchronous interaction, but they provide proven building blocks once a capability has been identified as a candidate for asynchronous design.

Conclusion

Choosing between synchronous and asynchronous design is a business-outcome decision expressed through technical pattern selection. Applying this framework consistently across capability designs reduces ad hoc coupling, improves resiliency, and ensures the interaction style matches the true needs of the business process it supports.